

PREPRINT September 6, 2013

LSSC

Yohann Salaun¹ & Marc Lebrun²

¹ Polytechnique, France (yohann.salaun@polytechnique.org)

² CMLA, ENS Cachan, France (marc.lebrun@cmla.ens-cachan.fr)

Abstract

1 Overview

2 Theoretical Description

2.1 Non Local Means

Self similarities inside pictures have been studied for a long time. Efros and Leung used it for texture synthesis [1] and were followed by Buades, Coll and Morel for denoising purpose [2].

Considering a noisy picture \mathbf{y} of n pixels seen as a column vector in \mathbb{R}^n , \mathbf{y} is divided into n overlapping patches of equal size m . The i -th pixel of \mathbf{y} is noted $\mathbf{y}[i]$ and the patch centered in $\mathbf{y}[i]$ and of size m is noted \mathbf{y}_i . Then, when two different patches \mathbf{y}_i and \mathbf{y}_j have similar values, their denoised version should also be similar. Moreover, assuming that the noise follows a Gaussian distribution, averaging similar patches should destroy the noise information and thus denoise the patch. Thus the denoised pixel $\mathbf{x}[i]$ is obtained with a linear combination of the others pixel in the noisy picture \mathbf{y} weighted by the similarity of their corresponding patches:

$$\mathbf{x}[i] = \frac{1}{N} \sum_{j=1}^n G(\mathbf{y}_i - \mathbf{y}_j) \mathbf{y}[j] \quad (1)$$

where N is the normalization factor and G a Gaussian that weights the patch similarities.

Following this idea, Mairal et al. [5] divided the pictures into n overlapping patches in order to denoise through self-similarity methods.

2.2 Learned Sparse Coding

The second idea is to assume that the denoised patches can be approximated by a sparse linear combinations of elements from a basis set.

This basis set, that contains the denoised picture information, is called a dictionary $\mathbf{D} \in \mathbb{R}^{m \times k}$. The linear combination is represented by a vector $\boldsymbol{\alpha} \in \mathbb{R}^k$ called a code. The aim is to have a dictionary that is not redundant (independant columns), not too large ($k \ll n$) but that contains the whole picture information.

The denoising problem consists then of finding one dictionary for the picture and one sparse code for each patch. Thus it can be reformulated into a minimization problem where we look for the optimal dictionary and codes that are the most similar to each noisy patches:

$$\forall i \in [1; n], \min_{\alpha \in \mathbb{R}^k, \mathbf{D} \in \mathbb{R}^{m \times k}} \sum_{i=1}^n \|\alpha_i\|_p \text{ s.t. } \|\mathbf{y}_i - \mathbf{D}\alpha_i\|_2^2 \leq \epsilon \quad (2)$$

$\mathbf{D}\alpha_i$ is the estimate of the i^{th} denoised patch which should be similar to the noisy patch and ϵ can be chosen according to the value of the estimated standard deviation of the noise.

Once the dictionary and the codes are learned, for each pixel, we have m estimations (from the m overlapping patches that contain it) and their averaging give us the denoised information of this pixel:

$$\forall i \in [1, n], \mathbf{x}[i] = \frac{1}{m} \sum_{j=1}^m \mathbf{D}\alpha_{\sigma(i,j)} \quad (3)$$

Where $\sigma(i, j)$ is the number of the patch that put pixel $\mathbf{x}[i]$ in j^{th} position.

Equation (2) is usually minimized with $p = 0$ or 1 . It becomes NP-hard to solve when $p = 0$ but a greedy algorithm such as Orthogonal Matching Pursuit [7] can quickly give a good approximation. The problem is convex with $p = 1$ and is efficiently solved with the Least Angle Regression algorithm [3]. Experimental observations [8] have shown that the learning part is better with $p = 1$ and the recomposition part with $p = 0$.

2.2.1 Learning Dictionary

In this part, we consider that $p = 1$ in equation (2). The problems becomes convex and the equation is only solved in order to find the optimal dictionary \mathbf{D} that could represents the denoised information of picture \mathbf{y} .

The method consists in updating the dictionary iteratively with only a small distribution of T patches. However they are chosen so that they are independently and identically distributed in the picture. First, the initial dictionary is learned offline on the 10 000 images of the PASCAL VOC'07 database using the online dictionary learning procedure of [4]. This procedure is then used on the noisy picture \mathbf{y} in order to improve the dictionary efficiency.

For each patch \mathbf{y}_i , the corresponding code is computed with the current dictionary kept constant and then the dictionary is updated with all the previous codes.

Thus, we minimize iteratively the two equations below:

- For $\mathbf{D} \in \mathbb{R}^{m \times k}$ and patch t , $\alpha^t = \operatorname{argmin}_{\alpha \in \mathbb{R}^k} \|\alpha\|_1$ s.t. $\|\mathbf{y}_t - \mathbf{D}^{t-1}\alpha\|_2^2 \leq \lambda$
- For patches 1 to t and their codes $\alpha \in \mathbb{R}^k$, $\mathbf{D}^t = \operatorname{argmin}_{\mathbf{D} \in \mathbb{R}^{m \times k}} \frac{1}{t} \sum_{i=1}^t (\frac{1}{2} \|\mathbf{y}_i - \mathbf{D}\alpha^i\|_2^2 + \lambda \|\alpha^i\|_1)$

The second equation can be re-written:

$$\mathbf{D}^t = \operatorname{argmin}_{\mathbf{D} \in \mathbb{R}^{m \times k}} \frac{1}{t} \left(\frac{1}{2} \operatorname{Tr}(\mathbf{D}^T \mathbf{D} \mathbf{B}^t) - \operatorname{Tr}(\mathbf{D}^T \mathbf{C}^t) \right) \quad (4)$$

where \mathbf{B} and \mathbf{C} are incrementally updated with:

- $\mathbf{B}^t \leftarrow \mathbf{B}^{t-1} + \alpha^t \alpha^{tT}$
- $\mathbf{C}^t \leftarrow \mathbf{C}^{t-1} + \mathbf{y}^t \alpha^{tT}$

The dictionary update is thus computed through a block-coordinate descent with warmrestarts [9]. This is an iterative method that updates the columns of the dictionary needing only matrices B and C and no matrix inversion in opposite to other approaches as Newton method. Moreover, since the convex optimization problem (Equation (4)) admits separable constraints in the updated blocks (columns), convergence to a global optimum is guaranteed. However Mairal et al empirically found that a single iteration of the dictionary update was sufficient to achieve convergence.

2.2.2 Denoising

TODO: MARC :)

2.3 Simultaneous Sparse Coding

2.3.1 Clustering

2.3.2 Simultaneous Least Angle Regression

2.3.3 Simultaneous Orthogonal Matching Pursuit

3 Notations

4 Algorithm Description

4.1 Algorithm Overview

The first part consists in initializing a dictionary that will denoise roughly the picture.

Once the picture is denoised a first time, a clustering is made in order to regroup similar patches for further treatment.

Then, iteratively for each clusters, the dictionary is updated using simultaneous sparse coding and the cluster is denoised.

4.2 Dictionnary Initialization

Input : number of iterations \mathbf{T} , i.i.d. sampling of \mathbf{T} patches of the noisy picture \mathbf{Y}_T , initial dictionary $\mathbf{D}^0 \in \mathbb{R}^{m \times k}$, regularization parameter λ

Output : learned dictionary \mathbf{D}

Initialization : $\mathbf{A}^0 \in \mathbb{R}^{k \times k} \leftarrow 0$, $\mathbf{B}^0 \in \mathbb{R}^{m \times k} \leftarrow 0$

for $t = 1..T$ **do**

$\mathbf{y}_t = \mathbf{Y}_T[t]$

 Sparse coding: compute with LARS algorithm:

$$\boldsymbol{\alpha}^t = \underset{\boldsymbol{\alpha} \in \mathbb{R}^k}{\operatorname{argmin}} \|\boldsymbol{\alpha}\|_1 \text{ s.t. } \|\mathbf{y}_t - \mathbf{D}^{t-1} \boldsymbol{\alpha}\|_2^2 \leq \lambda$$

$\mathbf{A}^t \leftarrow \mathbf{A}^{t-1} + \boldsymbol{\alpha}^t \boldsymbol{\alpha}^{tT}$

$\mathbf{B}^t \leftarrow \mathbf{B}^{t-1} + \mathbf{y}_t \boldsymbol{\alpha}^{tT}$

 Update dictionary from \mathbf{D}^{t-1} to \mathbf{D}^t so that:

$$\mathbf{D}^t = \underset{\mathbf{D} \in \mathbb{R}^{m \times k}}{\operatorname{argmin}} \frac{1}{t} \sum_{i=1}^t \left(\frac{1}{2} \|\mathbf{y}_t^i - \mathbf{D} \boldsymbol{\alpha}^i\|_2^2 + \lambda \|\boldsymbol{\alpha}^i\|_1 \right)$$

end

return \mathbf{D}^T

Algorithm 1: Online Dictionary Learning

Input :input dictionary $\mathbf{D} = [\mathbf{d}^1, \dots, \mathbf{d}^k] \in \mathbb{R}^{m \times k}$,

$\mathbf{A} = [\mathbf{a}^1, \dots, \mathbf{a}^k] \in \mathbb{R}^{k \times k}$, $\mathbf{B} = [\mathbf{b}^1, \dots, \mathbf{b}^k] \in \mathbb{R}^{m \times k}$

Output : updated dictionary \mathbf{D}

repeat

for $j = 1..k$ **do**

 update the j^{th} column:

if $\mathbf{A}(j,j) = 0$ **then**

$$\mathbf{d}^j \leftarrow 0$$

end

else

$$\mathbf{u}^j \leftarrow \frac{1}{\mathbf{A}(j,j)} (\mathbf{b}^j - \mathbf{D} \mathbf{a}^j) + \mathbf{d}^j$$

$$\mathbf{d}^j \leftarrow \frac{1}{\max(\|\mathbf{u}^j\|_2, 1)} \mathbf{u}^j$$

end

end

until convergence *Marc: apparemment, dans le code on a une boucle sur*

params.updateIteration = 1. Est-ce qu'il ne vaudrait mieux pas calculer l'argmin à chaque

boucle et s'arrêter lorsque c'est plus petit qu'une certaine valeur ?Yohann: je n'ai pas trop

compris ce que tu voulais faire. Ca c'est le pseudo code prsent par Mairal, qui est donc une

sorte de descente de gradient. Ce genre d'algo converge en thorie l'infini, en pratique on

prend un grand nombre d'itration. Cependant, dans ce cas on fait une minimisation alterne

d'un problme plus global (selon D puis selon alpha et on itre). Du coup, et surement cause de

contraintes de temps, Mairal a fix le nombre d'itration 1 mais permet de le changer en

paramtre dans son algo. Du coup j'ai voulu faire pareil que Mairal. ;

return \mathbf{D}

Algorithm 2: Dictionary Update *updateDictionary* *Marc: Algo validé*

Marc: TODO : choisir entre les indices ou les crochets pour les indices

Input : Input dictionary $\mathbf{D} \in \mathbb{R}^{m \times k}$, noisy patch $\mathbf{y} \in \mathbb{R}^m$, constraint $\lambda \in \mathbb{R}$

Output : code $\alpha \in \mathbb{R}^k$

—INITIALIZATION—

$\mathbf{G} \in \mathbb{R}^{k \times k} \leftarrow \mathbf{D}^T \mathbf{D}$

$\text{normPatch} \in \mathbb{R}^+ \leftarrow \|\mathbf{y}\|_2^2$

$\alpha \in \mathbb{R}^k \leftarrow \mathbf{0}$

$\mathcal{A} \in [0; k]^k \leftarrow \mathbf{0}$

Most correlated element

$\hat{\mathbf{c}} \in \mathbb{R}^k \leftarrow \mathbf{D}^T \mathbf{y}$

$C \in \mathbb{R}^+ \leftarrow \max_{j=1..k} (|\hat{\mathbf{c}}_j|)$

$\text{currentInd} \leftarrow j \text{ s.t. } \hat{\mathbf{c}}_j = C$

$\text{newAtom} \leftarrow \mathbf{True}$

if $\text{normPatch} < \lambda$ **then**

 | **return** $\mathbf{0}$

end

for $i = 1..k$ **do**

 | LOOP, see below

end

return α

Algorithm 3: LARS algorithm - Mairal Version `computeLars`

Marc: TODO : choisir entre les indices ou les crochets pour les indices

—NEW ATOM—

if *newAtom* **then**

$\mathcal{A}[i] \leftarrow \text{currentInd}$
 $G_A \in \mathbb{R}^{k \times i}, G_S \in \mathbb{R}^{i \times i}$
 $G_A[i^{th} \text{ column}] \leftarrow G[\text{currentInd}^{th} \text{ column}]$
 $G_S[i^{th} \text{ line}] \leftarrow G_A[\text{currentInd}^{th} \text{ line}]$
 symmetrize G_S
 UPDATE G_S^{-1}

end

—VARIABLES UPDATES—

$u \in \mathbb{R}^i \leftarrow G_S^{-1}(\text{sgn}(\hat{c}_{\mathcal{A}[j]}))_{j \in [1;i]}$

$C \leftarrow |\hat{c}[\mathcal{A}[1]]| = \max_{j=1..k}(|\hat{c}_j|)$

$\gamma \in \mathbb{R}_*^+ \leftarrow \min^+ \left(\frac{C + \hat{c}_j}{1 + (G_A u)[j]}, \frac{C - \hat{c}_j}{1 - (G_A u)[j]} \right)_{j \text{ s.t. } \mathcal{A}[j]=0}$

$\text{currentInd} = j \text{ s.t. } \gamma = \frac{C \pm \hat{c}_j}{1 \pm (G_A u)[j]}$

$\text{ratio} \in \mathbb{R}^i \leftarrow \left(-\frac{\alpha[\mathcal{A}[j]]}{u_j} \right)_{j \in [1;i]}$

$\text{stepDownDate} \in \mathbb{R}_*^+ \leftarrow \min^+(\text{ratio})$ (\min^+ is the minimum between strictly positive values only)

$\text{downDateInd} \in [1;k] \leftarrow \mathcal{A}[j] \text{ s.t. } \text{ratio}[j] = \text{stepDownDate}$

—POLYNOMIAL RESOLUTION—

$a \in \mathbb{R} \leftarrow \sum_{j \in [1;i]} \text{sgn}(\hat{c}[\mathcal{A}[j]])u[j]$

$b \in \mathbb{R} \leftarrow \sum_{j \in [1;i]} \hat{c}[\mathcal{A}[j]]u[j]$

$c \in \mathbb{R} \leftarrow \text{normPatch} - \lambda$

$\Delta \in \mathbb{R} \leftarrow b^2 - ac$

$\text{stepMAX} \in \mathbb{R}^+ \leftarrow \min(\frac{b - \sqrt{\Delta}}{a}, C)$

—FINAL STEP & BREAK—

$\gamma \leftarrow \min(\gamma, \text{stepDownDate}, \text{stepMAX})$

for $j = 1..i$ **do**

$\alpha[\mathcal{A}[j]] \leftarrow \alpha[\mathcal{A}[j]] + \gamma u[j]$

end

$\hat{c} \leftarrow \hat{c} - \gamma G_A u$

$\text{normPatch} \leftarrow \text{normPatch} + a\gamma^2 - 2b\gamma$

if $|\gamma| < 10^{-6}$ **or** $\gamma = \text{stepMAX}$ **or** $\text{normPatch} < 10^{-6}$ **or** $\text{normPatch} - \lambda < 10^{-6}$ **then**
 break

end

if $\gamma = \text{stepDownDate}$ **then**

 DOWNDATA G_S^{-1} w.r.t downDateInd

$\mathcal{A}[\text{downDateInd}] \leftarrow 0$

$\alpha[\text{downDateInd}] \leftarrow 0$

 newAtom \leftarrow **False**

$i \leftarrow i - 1$

end

else

 newAtom \leftarrow **True**

$i \leftarrow i + 1$

end

Algorithm 4: LARS algorithm - Mairal Version **computeLars** Marc: Algo correspondant, mais à valider. - LOOP

Input : Gram matrix $G_S \in \mathbb{R}^{i \times i}$, and its former inverse to update $G_S^{-1} \in \mathbb{R}^{i-1 \times i-1}$

Output : updated $G_S^{-1} \in \mathbb{R}^{i \times i}$

if $i = 1$ **then**

return $\frac{1}{G_S}$

end

$u \leftarrow G_S^{-1} G_S[i^{th} \text{ line}]$

$\sigma \leftarrow \frac{1}{G_S(i,i) - u \cdot G_S[i^{th} \text{ line}]}$

$G_S^{-1}(i, i) \leftarrow \sigma$

$G_S^{-1}[i^{th} \text{ line}] \leftarrow -\sigma u$

return $G_S^{-1} \leftarrow G_S^{-1} + \sigma u u^T$

Algorithm 5: Update invert algorithm **updateGram** Marc: Algo correspondant, mais à valider.

Input : pseudo-Gram matrix $G_A \in \mathbb{R}^{k \times i}$ Gram matrix $G_S \in \mathbb{R}^{i \times i}$, and its inverse $G_S^{-1} \in \mathbb{R}^{i \times i}$, criticalInd $\in [1; k]$, current iteration i

Output : downdated matrices $G_A \in \mathbb{R}^{k \times i-1}$, $G_S, G_S^{-1} \in \mathbb{R}^{i-1 \times i-1}$

$\sigma \leftarrow \frac{1}{G_S^{-1}(\text{criticalInd}, \text{criticalInd})}$

$u \leftarrow G_S^{-1}[\text{criticalInd}^{th} \text{ line}]$ without its criticalIndth coefficient

for $j = \text{criticalInd} : i-1$ **do**

$G_A[j^{th} \text{ column}] \leftarrow G_A[(j+1)^{th} \text{ column}]$

for $k = 1 : \text{criticalInd}-1$ **do**

$G_S(j, k) \leftarrow G_S(j+1, k)$

$G_S^{-1}(j, k) \leftarrow G_S^{-1}(j+1, k)$

end

for $k = \text{criticalInd} : i$ **do**

$G_S(j, k) \leftarrow G_S(j+1, k+1)$

$G_S^{-1}(j, k) \leftarrow G_S^{-1}(j+1, k+1)$

end

end

$G_S^{-1} \leftarrow G_S^{-1} - \sigma u u^T$

Algorithm 6: Downdate invert algorithm **downdateGram** Marc: Algo correspondant, mais à valider.

Glossary

Image Credits



© IPOL (there's no need to credit this image, here is used as an example.)

5 References

References

- [1] A. Efros, and T. Leung *Texture synthesis by non-parametric sampling*. ICCV, 1999.
- [2] A. Buades, B. Coll, and J. Morel *A non-local algorithm for image denoising*. CVPR, 2005.
- [3] B. Efron, T. Hastie, I. Johnstone, and R. Tibshirani *Least angle regression*. Ann. Statist., 32(2):407499, 2004.

- [4] J. Mairal, F. Bach, J. Ponce, and G. Sapiro *Online dictionary learning for sparse coding*. ICML, 2009.
- [5] J. Mairal, F. Bach, J. Ponce, G. Sapiro and A. Zisserman *Non-local Sparse Models for Image Restoration*. International Conference on Computer Vision, 2009.
- [6] J. Mairal *Représentations parcimonieuses en apprentissage statistique, traitement d'image et vision par ordinateur*. PhD thesis, 2010.
- [7] S. Mallat and Z. Zhang *Matching pursuit in a timefrequency dictionary*. IEEE T. SP, 41(12):3397-3415, 1993.
- [8] M. Elad and M. Aharon *Image denoising via sparse and redundant representations over learned dictionaries*. IEEE T. IP, 54(12):3736-3745, 2006.
- [9] D. P. Bertsekas *Nonlinear programming*. Athena Scientific Belmont, 1999.